Decenter: Worry Map Journal PRD

Draft v1 · 29 Oct 2026 · Built from the clinical PRD plus our product discussion

1. Summary

Decenter is a diary-style journal where each worry gets its own page, and that page is a map that grows from left to right. The user writes a worry, answers plain clarifying questions, and only then reaches a conclusion: a defused version of the thought, a split between what they can and cannot steer, and a small plan. The AI helps structure the page and synthesize what the user wrote. It does not answer the worry.

Design reference: the two-screen canvas at https://claude.ai/artifact/1Ln7SEwYC7ApbSFZCm6v2t (Worry page and Conclusion).

2. The problem

2.1 The user problem

People who worry chronically get stuck in loops. Most journaling and self-help tools make this worse because they work on the content of the worry: debating it, logging it, listing it. That teaches the user that thoughts need to be fought. The clinical PRD frames this through metacognitive therapy and acceptance and commitment therapy:

- Type 1 worry is about an external threat or a bodily sensation. Type 2 worry is worry about the worry itself ("I can't control this"), and it is what turns a passing concern into something heavier.
- Worry is mostly abstract and verbal, which avoids the concrete, sensory details that let the emotion settle. Vague language ("everything is ruined") keeps the loop going.
- Users operate in object mode, where a thought feels like a fact. The goal is metacognitive mode, where the thought is seen as a passing mental event.
- Binary worry trees ("can you control it, yes or no") overload the user. Real situations are mixed, so control needs to be a spectrum.

2.2 The product problem

The first prototype, generated in Google AI Studio, got the sequence right but the grammar wrong. Every step was a large question card, the canvas had to zoom out to 35 percent to fit, and the questions read like story prompts instead of real clarifying questions. The tool could not produce the page layout we want.

3. The solution

A worry page is a tree, and the tree is the diary entry.

- Questions sit on the lines between nodes, not in big cards.
- Boxes appear only where the user writes.
- Brackets merge the answers into an output the AI drafts from the user's own words.
- The conclusion is locked until the facts are written down and confirmed, so the user cannot skip the concrete step.

4. Decisions made so far

- **Structure:** a fixed trunk taken from the clinical PRD (unpack, defuse, control spectrum, action). The AI writes the question wording and the follow-up branches, and the user can add their own questions.
- **AI role:** the AI synthesizes what the user wrote. It does not answer or reassure.
- **Page model:** each worry is a diary page, and the map is the entry itself, so past entries can be browsed like a diary.
- **Control dial:** each part of a worry carries its own dial from 0 (cannot steer) to 1 (can steer). Below 0.5 routes to acceptance. 0.5 and above routes to an action plan.
- **Question style:** concrete, investigative questions (what happened, when was it, who was there, how did it feel) asked before any conclusion. Not story prompts.
- **v1 safety layer:** Quick Calm and a session timer only.

5. User flow

1. **Start a page.** The user taps New page and writes the worry in their own words.
2. **Answer the clarifying questions.** The first-level questions are fixed. Each answer can open an AI follow-up branch (for example, "What was said?" or "Where in your body?"). The user can also add their own question.

3. **Confirm the facts card.** The AI drafts a short, factual summary from the answers. The user edits it and confirms.
4. **Conclude.** The conclusion unlocks. The user taps "I'm ready to conclude".
5. **Defuse the thought.** The thought is reframed as "I'm having the thought that..." and the user marks it as a passing mental event or a present fact.
6. **Split by control.** The user breaks the worry into parts and sets a dial on each.
7. **Plan.** Parts under 0.5 go to "Make room for". Parts at 0.5 and above go to "If, then, until".
8. **Park the rest.** The user picks a 15-minute worry window and returns to the present.

6. Functional requirements

6.1 Diary

- A sidebar lists past pages with date, category and a short title.
- New page starts a blank tree.
- Each page opens to its saved tree and conclusion.

6.2 Worry page (the tree)

- The worry text is the root node on the left.
- First-level questions: What happened? When was it? Who was there? How did it feel? Anything else?
- Each question is a label on its branch. Each answer is a writing box at the end of the branch.
- AI follow-up branches are visually distinct (dashed, purple) from the user's own boxes.
- A branch with no follow-up simply connects through to the bracket.
- A "+ Add your own question" slot lets the user add a branch.
- All answers merge through a bracket into the facts card.

6.3 Facts card

- AI drafts it using only the user's words.
- The user can edit it. Confirming it unlocks the conclusion.
- If the user's answers are still abstract (words like always, never, ruined, disaster), the app asks for a concrete detail before the card can be drafted.

6.4 Ready gate

- Proposed: the user decides when to conclude by tapping "I'm ready to conclude". After about three answered questions the AI offers a soft nudge. The AI does not block the user once the facts card is confirmed.

6.5 Conclusion

- **Thought card:** the worry rewritten as "I'm having the thought that..." with a choice between "A fact, right now" and "A passing mental event". If the user chooses fact, show a short prompt about who is observing the thought, then ask again.
- **Parts and dials:** the user lists parts of the worry and sets a 0.00 to 1.00 dial on each.
- **Make room for:** a willingness statement in the form "I'm willing to make room for [feeling] while I [valued activity]."
- **If, then, until:** three required fields. The stop signal must be an external milestone (a time, a finished item), not a feeling.
- **Worry window:** a single 15-minute slot later in the day, with a commitment card.

6.6 Quick Calm and session timer

- Quick Calm is always visible. It opens a low-stimulation screen with a short grounding exercise.
- A session timer shows time on the current page (20 minutes is the working limit). It informs the user. It does not lock the page in v1.

7. AI behavior rules

- Writes question wording within the fixed slots and suggests follow-ups from the user's answers.
- Drafts the facts card, thought card and plan summaries only from what the user wrote.
- Never reassures ("this will be fine"), diagnoses, or argues with the content of the worry.
- Never invents facts about the situation.
- Every AI output is editable and needs user confirmation.

8. Data fields (from the clinical PRD)

- raw_worry_text (string, max 1000 characters)
- timestamp_anchor (date and time of the triggering event)
- facts list (answers to the clarifying questions, plus the confirmed facts card)
- parsed_thought_string and defusion frame
- control_spectrum_slider_value (0.00 to 1.00) and category (below 0.50 uncontrollable, 0.50 and above controllable)
- selected_value, if_trigger_condition, then_action_commitment, execution_stop_signal_criteria (each up to 200 characters)
- worry_window_timestamp

9. Design direction

- Desktop first. The canvas shows a 1280 px page with a diary sidebar and a left-to-right tree.
- Calm, paper-like look with serif headings. Sage green marks the user's own input. Purple marks AI output.
- Questions are small labels on lines. Nothing in the tree is a large card except the user's own boxes and the AI outputs.
- Quick Calm is a soft red-tinted button in the sidebar. It should be easy to find and never alarming.

10. Out of scope for v1

These safeguards are in the clinical PRD and are deliberately deferred:

- Panic and hyperarousal detection (keywords, typing speed)
- The rumination circuit-breaker that locks fields
- Daily logging caps and compulsive-checking limits
- The severe meta-worry overlay and uncontrollability experiment
- Mobile layout

11. Success measures

The app should not reward frequent use, so engagement is the wrong goal. Better signals:

- Users reach the conclusion on pages they start.
- Belief in the thought (the 1 to 5 rating after defusion) drops between start and end of a page.
- Time per page stays under the 20-minute working limit.
- Users complete the "If, then, until" plan with an external stop signal.
- Users do not open many pages in a single day.

12. Open questions

- Who decides when the facts are enough: the user alone, or a soft AI nudge too?
- How many follow-up branches per answer, and how deep, before the tree gets heavy?
- Does the AI suggest the parts of the worry for the control step, or does the user list them from scratch?
- What should happen if a user writes something that signals serious distress? Even without the full safeguard layer, v1 likely needs a visible path to human help.
- What is the content of Quick Calm in v1?
- Mobile: does the tree collapse into a vertical list, or stay a map?

13. Next steps

- Make the worry page interactive so a follow-up branch appears after an answer.
- Design the Quick Calm overlay.
- Design the mobile version.
- Write the question prompts and AI guardrail prompts.